

KYB Tool - Technical Architecture Document

Part 1: System Design, Microservices, and Data Architecture

Document Information

- **Document Type:** Technical Architecture Document
 - **Project:** KYB Tool - Enterprise-Grade Know Your Business Platform
 - **Version:** 1.0
 - **Date:** January 2025
 - **Status:** Final Specification
 - **Pages:** Part 1 of 3
-

1. System Design Overview

1.1 Architecture Principles

Microservices-First Design

- Each service owns its data and business logic
- Services communicate via well-defined APIs and events
- Independent deployment and scaling capabilities
- Fault isolation and resilience patterns

Event-Driven Architecture

- Asynchronous processing for non-blocking operations
- Event sourcing for complete audit trails

- Real-time data streaming and processing
- Decoupled service communication


Cloud-Native Design


- Container-first deployment strategy
- Kubernetes orchestration for scalability
- Infrastructure as Code (IaC) for consistency
- Multi-region deployment for global scale

Security-by-Design

- Zero-trust network architecture
- End-to-end encryption for all data flows
- Comprehensive audit logging and monitoring
- Compliance-ready security controls

1.2 High-Level System Architecture





1.3 Technology Stack Selection

Backend Services

- **Go:** High-performance services (API Gateway, Data Ingestion, Monitoring)
- **Python:** AI/ML services (Classification, Risk Assessment, Analytics)
- **Rationale:** Go for I/O intensive operations, Python for ML workloads

Data Layer

- **PostgreSQL:** Primary transactional database with ACID guarantees
- **Redis Cluster:** Distributed caching and session management
- **Elasticsearch:** Full-text search and log aggregation
- **Apache Kafka:** Event streaming and message processing

Infrastructure

- **Kubernetes**: Container orchestration and scaling
 - **Istio**: Service mesh for communication and security
 - **Prometheus/Grafana**: Monitoring and observability
 - **ELK Stack**: Centralized logging and analytics
-

2. Microservices Architecture

2.1 Service Decomposition Strategy

Domain-Driven Design Principles

- Services aligned with business capabilities
- Clear ownership and responsibility boundaries
- Minimal cross-service data dependencies
- Independent evolution and deployment

Service Sizing Guidelines

- 2-pizza team ownership model
- Single responsibility principle
- Database per service pattern
- Independent scaling requirements

2.2 Core Services Definition

Authentication Service (Go)

```
yaml
```

Responsibilities:

- User authentication and authorization
- JWT token generation and validation
- API key management and rotation
- Multi-tenant access control
- Session management and SSO preparation

API Endpoints:

- POST /auth/login
- POST /auth/logout
- GET /auth/validate
- POST /auth/refresh
- CRUD /auth/api-keys

Database Schema:

- users (id, tenant_id, email, password_hash, role, created_at, updated_at)
- api_keys (id, tenant_id, key_hash, permissions, rate_limit, created_at, expires_at)
- sessions (id, user_id, token_hash, expires_at, last_accessed)

Dependencies:

- Redis for session storage
- PostgreSQL for user data

Classification Service (Python/FastAPI)

yaml

Responsibilities:

- Business classification using AI models
- MCC, NAICS, SIC code assignment
- Confidence scoring and alternatives
- Website content analysis
- Product/service categorization

API Endpoints:

- POST /classify/business
- POST /classify/batch
- GET /classify/{job_id}/status
- POST /classify/website
- GET /classify/codes/search

ML Models:

- BERT-based business description classifier
- XGBoost ensemble for multi-code prediction
- Custom similarity matching algorithms
- Website content categorization CNN

Data Sources:

- Business registration data
- Website content scraping
- Industry classification databases
- Product/service taxonomies

Risk Assessment Service (Python/FastAPI)

yaml

Responsibilities:

- Real-time risk scoring (1-100 scale)
- Predictive risk modeling (3, 6, 12 months)
- Fraud detection and pattern analysis
- Regulatory compliance risk assessment
- Risk factor explanation and recommendations

API Endpoints:

- POST /risk/assess
- POST /risk/predict
- GET /risk/{merchant_id}/history
- POST /risk/batch-assess
- GET /risk/factors/explain

ML Models:

- XGBoost + LSTM ensemble for risk prediction
- Isolation Forest for anomaly detection
- Graph Neural Networks for transaction analysis
- SHAP for model explainability

Risk Categories:

- Operational Risk (business stability, operational issues)
- Financial Risk (credit, cash flow, financial distress)
- Regulatory Risk (compliance violations, sanctions)
- Reputational Risk (negative news, customer complaints)
- Cybersecurity Risk (data breaches, security incidents)

Data Ingestion Service (Go)

```
yaml
```

Responsibilities:

- Web scraping and content extraction
- Third-party API integration management
- Data validation and normalization
- Rate limiting and quota management
- Error handling and retry logic

API Endpoints:

- POST /ingest/website
- POST /ingest/business-registry
- GET /ingest/job/{id}/status
- POST /ingest/validate
- GET /ingest/sources/health

Data Sources:

- Business registry APIs (Secretary of State databases)
- Website content scraping (Playwright-based)
- Social media APIs (when available and free)
- Government databases (OFAC, tax authorities)
- Open data sources (OpenCorporates, etc.)

Features:

- Intelligent bot evasion techniques
- Proxy rotation and IP management
- Content quality assessment
- Duplicate detection and deduplication

Compliance Service (Go)

yaml

Responsibilities:

- Regulatory compliance checking
- Sanctions screening (OFAC, UN, EU)
- Audit trail generation
- Compliance reporting
- Policy enforcement

API Endpoints:

- POST /compliance/screen
- GET /compliance/audit-trail
- POST /compliance/report
- GET /compliance/policies
- POST /compliance/validate

Compliance Frameworks:

- OFAC sanctions screening
- EU sanctions compliance
- PCI DSS data handling
- GDPR privacy protection
- SOC 2 security controls

Features:

- Fuzzy name matching for sanctions screening
- Automated compliance reporting
- Policy violation detection
- Audit log immutability (blockchain-backed)

2.3 Service Communication Patterns

Synchronous Communication

- Direct API calls for real-time operations

- Circuit breaker pattern for fault tolerance
- Timeout and retry policies
- Service mesh for secure communication

Asynchronous Communication

- Event-driven messaging via Kafka
- Event sourcing for audit trails
- Saga pattern for distributed transactions
- Dead letter queues for failed messages

Data Consistency Strategy

- Eventual consistency between services
 - Distributed transaction coordination when required
 - Idempotent operations for retry safety
 - Conflict resolution strategies
-

3. Database Design and Data Architecture

3.1 Data Architecture Overview

Multi-Store Strategy

Data Architecture

Operational Data (OLTP)

- └── PostgreSQL Primary Cluster
 - | └── Tenant Management
 - | └── User & Authentication
 - | └── Business Profiles
 - | └── Risk Assessments
 - | └── Compliance Records

Caching Layer

- └── Redis Cluster
 - | └── Session Management
 - | └── API Response Caching
 - | └── ML Model Cache
 - | └── Real-time Counters

Search & Analytics (OLAP)

- └── Elasticsearch Cluster
 - | └── Full-text Search
 - | └── Log Aggregation
 - | └── Real-time Analytics

Machine Learning

- └── Feature Store (Feast)
 - | └── Business Features
 - | └── Risk Features
 - | └── Model Training Data

Data Warehouse (Future)

- └── PostgreSQL Analytics DB

- |— Historical Data
- |— Reporting Data
- |— Business Intelligence

3.2 PostgreSQL Schema Design

Core Business Entities

sql

```
-- Tenant Management
CREATE TABLE tenants (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    name VARCHAR(255) NOT NULL,
    plan_type VARCHAR(50) NOT NULL,
    api_rate_limit INTEGER DEFAULT 100,
    features JSONB DEFAULT '{}',
    billing_info JSONB,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    deleted_at TIMESTAMP WITH TIME ZONE NULL
);
```

```
-- User Management
CREATE TABLE users (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    email VARCHAR(255) UNIQUE NOT NULL,
    password_hash VARCHAR(255) NOT NULL,
    role VARCHAR(50) NOT NULL DEFAULT 'user',
    permissions JSONB DEFAULT '[]',
    last_login TIMESTAMP WITH TIME ZONE,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    deleted_at TIMESTAMP WITH TIME ZONE NULL
);
```

```
-- Business Profiles
CREATE TABLE businesses (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    external_id VARCHAR(255), -- Customer's merchant ID
    legal_name VARCHAR(500) NOT NULL,
```

```
    dba_name VARCHAR(500),
    tax_id VARCHAR(50),
    business_type VARCHAR(100),
    registration_state VARCHAR(50),
    registration_country VARCHAR(50),
    registration_date DATE,
    website_url TEXT,
    business_description TEXT,
    contact_info JSONB,
    address JSONB,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    updated_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),

    -- Indexes for performance
    INDEX idx_businesses_tenant_external (tenant_id, external_id),
    INDEX idx_businesses_tax_id (tax_id) WHERE tax_id IS NOT NULL,
    INDEX idx_businesses_website (website_url) WHERE website_url IS NOT NULL
);

-- Business Classifications
CREATE TABLE business_classifications (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    business_id UUID NOT NULL REFERENCES businesses(id) ON DELETE CASCADE,
    classification_type VARCHAR(20) NOT NULL, -- 'MCC', 'NAICS', 'SIC'
    code VARCHAR(20) NOT NULL,
    description TEXT,
    confidence_score DECIMAL(5,4), -- 0.0000 to 1.0000
    is_primary BOOLEAN DEFAULT false,
    source VARCHAR(100), -- 'ai_model', 'manual', 'api'
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),

    INDEX idx_classifications_business (business_id),
    INDEX idx_classifications_type_code (classification_type, code)
```

```
);

-- Risk Assessments
CREATE TABLE risk_assessments (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    business_id UUID NOT NULL REFERENCES businesses(id) ON DELETE CASCADE,
    assessment_type VARCHAR(50) NOT NULL, -- 'initial', 'periodic', 'triggered'
    overall_score INTEGER NOT NULL CHECK (overall_score BETWEEN 1 AND 100),
    risk_level VARCHAR(20) NOT NULL, -- 'low', 'medium', 'high', 'critical'

    -- Individual risk scores
    operational_risk INTEGER CHECK (operational_risk BETWEEN 1 AND 100),
    financial_risk INTEGER CHECK (financial_risk BETWEEN 1 AND 100),
    regulatory_risk INTEGER CHECK (regulatory_risk BETWEEN 1 AND 100),
    reputational_risk INTEGER CHECK (reputational_risk BETWEEN 1 AND 100),
    cybersecurity_risk INTEGER CHECK (cybersecurity_risk BETWEEN 1 AND 100),

    risk_factors JSONB DEFAULT '[]',
    recommendations JSONB DEFAULT '[]',
    model_version VARCHAR(50),
    confidence_interval JSONB, -- {'lower': 0.8, 'upper': 0.95}

    -- Predictive scores
    risk_trend VARCHAR(20), -- 'increasing', 'stable', 'decreasing'
    predicted_3m INTEGER CHECK (predicted_3m BETWEEN 1 AND 100),
    predicted_6m INTEGER CHECK (predicted_6m BETWEEN 1 AND 100),
    predicted_12m INTEGER CHECK (predicted_12m BETWEEN 1 AND 100),

    assessed_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    valid_until TIMESTAMP WITH TIME ZONE,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),

    INDEX idx_risk_assessments_business (business_id),
```

```
INDEX idx_risk_assessments_score (overall_score),
INDEX idx_risk_assessments_date (assessed_at)
);

-- Website Analysis
CREATE TABLE website_analyses (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    business_id UUID NOT NULL REFERENCES businesses(id) ON DELETE CASCADE,
    url TEXT NOT NULL,
    content_hash VARCHAR(64), -- SHA-256 of content

    -- Content analysis results
    products_services JSONB DEFAULT '[]',
    prohibited_content JSONB DEFAULT '[]',
    risk_indicators JSONB DEFAULT '[]',
    content_quality_score DECIMAL(5,4),
    ssl_certificate JSONB,

    -- Technical analysis
    page_load_time INTEGER, -- milliseconds
    mobile_friendly BOOLEAN,
    seo_score INTEGER CHECK (seo_score BETWEEN 0 AND 100),

    -- Screenshots and evidence
    screenshot_url TEXT,
    evidence JSONB DEFAULT '{}',

    analyzed_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),

    INDEX idx_website_analyses_business (business_id),
    INDEX idx_website_analyses_url_hash (url, content_hash)
);
```

Compliance and Audit Tables

sql

```
-- Sanctions Screening
CREATE TABLE sanctions_screenings (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    business_id UUID NOT NULL REFERENCES businesses(id) ON DELETE CASCADE,
    screening_type VARCHAR(50) NOT NULL, -- 'ofac', 'un', 'eu', 'uk'
    entity_name VARCHAR(500) NOT NULL,
    match_found BOOLEAN NOT NULL DEFAULT false,
    match_details JSONB DEFAULT '[]',
    match_score DECIMAL(5,4),
    list_version VARCHAR(100),
    screened_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    INDEX idx_sanctions_business (business_id),
    INDEX idx_sanctions_match (match_found),
    INDEX idx_sanctions_date (screened_at)
);
```

```
-- Audit Trail
CREATE TABLE audit_logs (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    user_id UUID REFERENCES users(id),
    api_key_id UUID, -- For API-driven actions

    event_type VARCHAR(100) NOT NULL, -- 'business_created', 'risk_assessed', etc.
    entity_type VARCHAR(50) NOT NULL, -- 'business', 'user', 'assessment'
    entity_id UUID,

    action VARCHAR(50) NOT NULL, -- 'create', 'update', 'delete', 'read'
    old_values JSONB,
    new_values JSONB,

    ip_address INET,
```

```

    user_agent TEXT,
    request_id UUID,
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    INDEX idx_audit_tenant_date (tenant_id, created_at),
    INDEX idx_audit_entity (entity_type, entity_id),
    INDEX idx_audit_user (user_id)
);

-- API Usage Tracking
CREATE TABLE api_usage (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    tenant_id UUID NOT NULL REFERENCES tenants(id),
    api_key_id UUID,
    endpoint VARCHAR(200) NOT NULL,
    method VARCHAR(10) NOT NULL,
    status_code INTEGER NOT NULL,
    response_time INTEGER, -- milliseconds
    request_size INTEGER, -- bytes
    response_size INTEGER, -- bytes
    created_at TIMESTAMP WITH TIME ZONE DEFAULT NOW(),
    INDEX idx_api_usage_tenant_date (tenant_id, created_at),
    INDEX idx_api_usage_endpoint (endpoint)
);

```

3.3 Redis Caching Strategy

Cache Structure and TTL Policies

yaml

Session Management:

Key Pattern: "session:{token_hash}"

TTL: 24 hours (configurable)

Data: User session information, permissions

API Response Caching:

Key Pattern: "api:{endpoint}:{hash(params)}"

TTL: 1-24 hours (based on data freshness requirements)

Data: Serialized API responses

Business Classification Cache:

Key Pattern: "classify:{hash(business_data)}"

TTL: 7 days

Data: Classification results with confidence scores

Risk Assessment Cache:

Key Pattern: "risk:{business_id}:{assessment_date}"

TTL: 24 hours

Data: Risk scores and predictions

Website Analysis Cache:

Key Pattern: "website:{url_hash}:{content_hash}"

TTL: 7 days

Data: Website analysis results

Rate Limiting:

Key Pattern: "rate_limit:{api_key_id}:{window}"

TTL: Based on rate limit window (1 minute, 1 hour, 1 day)

Data: Request counters

Feature Vectors (ML):

Key Pattern: "features:{business_id}:{feature_set}"

TTL: 1 hour

Data: Preprocessed ML feature vectors

Cache Warming and Invalidations

```
python

# Intelligent cache warming strategy
class CacheWarmer:
    async def warm_popular_classifications(self):
            """Pre-populate cache with popular business types"""
            popular_queries = await self.get_popular_classification_queries()
            for query in popular_queries:
                    await self.classification_service.classify(query, cache_only=False)

    async def warm_risk_models(self):
            """Pre-load ML models into memory/cache"""
            await self.risk_service.load_models()

    async def invalidate_on_update(self, business_id: UUID):
            """Invalidate related cache entries when business data updates"""
        patterns = [
            f"risk:{business_id}:*",
            f"classify:{*}{business_id}*",
            f"website:{*}{business_id}*" ]
        for pattern in patterns:
            await self.redis.delete_pattern(pattern)
```

3.4 Elasticsearch Schema and Usage

Index Templates for Different Data Types

json

```
{  
  "business_profiles": {  
    "mappings": {  
      "properties": {  
        "business_id": {"type": "keyword"},  
        "tenant_id": {"type": "keyword"},  
        "legal_name": {  
          "type": "text",  
          "analyzer": "standard",  
          "fields": {  
            "keyword": {"type": "keyword"}  
          }  
        },  
        "business_description": {  
          "type": "text",  
          "analyzer": "english"  
        },  
        "classifications": {  
          "type": "nested",  
          "properties": {  
            "type": {"type": "keyword"},  
            "code": {"type": "keyword"},  
            "description": {"type": "text"},  
            "confidence_score": {"type": "float"}  
          }  
        },  
        "risk_score": {"type": "integer"},  
        "risk_level": {"type": "keyword"},  
        "website_content": {"type": "text"},  
        "created_at": {"type": "date"},  
        "@timestamp": {"type": "date"}  
      }  
    }  
  }
```

```
        },  
  
        "audit_logs": {  
            "mappings": {  
                "properties": {  
                    "tenant_id": {"type": "keyword"},  
                    "user_id": {"type": "keyword"},  
                    "event_type": {"type": "keyword"},  
                    "entity_type": {"type": "keyword"},  
                    "action": {"type": "keyword"},  
                    "ip_address": {"type": "ip"},  
                    "details": {"type": "object"},  
                    "@timestamp": {"type": "date"}  
                }  
            }  
        },  
  
        "api_logs": {  
            "mappings": {  
                "properties": {  
                    "tenant_id": {"type": "keyword"},  
                    "endpoint": {"type": "keyword"},  
                    "method": {"type": "keyword"},  
                    "status_code": {"type": "integer"},  
                    "response_time": {"type": "integer"},  
                    "request_size": {"type": "integer"},  
                    "response_size": {"type": "integer"},  
                    "@timestamp": {"type": "date"}  
                }  
            }  
        }  
    }  
}
```

This completes Part 1 of the Technical Architecture document, covering system design, microservices architecture, and comprehensive data architecture. The document provides specific implementation details for database schemas, caching strategies, and search capabilities.

Should I proceed with **Part 2: AI/ML Architecture, External Integrations, and Security?**