

KYB Tool - Technical Architecture Document

Part 3: Performance, Infrastructure, and Operations

Document Information

- **Document Type:** Technical Architecture Document
 - **Project:** KYB Tool - Enterprise-Grade Know Your Business Platform
 - **Version:** 1.0
 - **Date:** January 2025
 - **Status:** Final Specification
 - **Pages:** Part 3 of 3
-

7. Performance Optimization Architecture

7.1 Caching Strategy and Implementation

Multi-Layer Caching Architecture

Caching Architecture

CDN Layer (Cloudflare)

- └── Static Assets (JS, CSS, Images)
- └── API Response Caching (GET endpoints)
- └── Geographic Edge Locations
- └── DDoS Protection & Bot Mitigation

|



Application Load Balancer Cache

- └── Session Affinity
- └── Health Check Results
- └── SSL Termination
- └── Request Routing Rules

|



API Gateway Cache (Kong)

- └── Rate Limiting Counters
- └── Authentication Token Cache
- └── Response Caching (5min-24hrs TTL)
- └── Request/Response Transformation

|



Application Layer Cache (Redis Cluster)

- └── Business Classification Cache (7 days TTL)
- └── Risk Assessment Cache (24 hours TTL)
- └── Website Analysis Cache (7 days TTL)
- └── ML Model Predictions (1 hour TTL)
- └── Session Data (24 hours TTL)

|



Database Query Cache (PostgreSQL)

 └─ Query Plan Cache

 └─ Prepared Statement Cache

 └─ Index Buffer Pool

 └─ Connection Pool

Intelligent Cache Management

python

```
import asyncio
import hashlib
import json
from typing import Any, Dict, Optional
from datetime import datetime, timedelta

class IntelligentCacheManager:
    """
    AI-powered cache management with predictive warming and smart invalidation
    """

    def __init__(self, redis_cluster, ml_predictor):
        self.redis = redis_cluster
        self.predictor = ml_predictor
        self.hit_rate_target = 0.95
        self.cache_strategies = {
            'classification': {'ttl': 7*24*3600, 'priority': 'high'},
            'risk_assessment': {'ttl': 24*3600, 'priority': 'high'},
            'website_analysis': {'ttl': 7*24*3600, 'priority': 'medium'},
            'business_registry': {'ttl': 30*24*3600, 'priority': 'medium'},
            'sanctions_screening': {'ttl': 24*3600, 'priority': 'high'}
        }

    async def get_with_intelligent_refresh(self, key: str, fetch_func, cache_type: str) -> Any:
        """
        Get cached value with intelligent refresh prediction
        """

        # Try to get from cache
        cached_data = await self.redis.get(key)

        if cached_data:
            data = json.loads(cached_data)
```

```
# Check if we should proactively refresh
if await self.should_proactive_refresh(key, data, cache_type):
    # Refresh in background
    asyncio.create_task(self.background_refresh(key, fetch_func, cache_type))

    return data['value']

# Cache miss - fetch and store
value = await fetch_func()
await self.set_with_metadata(key, value, cache_type)
return value

async def should_proactive_refresh(self, key: str, cached_data: Dict, cache_type: str) -> bool:
    """
    Predict if cache entry should be refreshed before expiration
    """

    # Get cache metadata
    cached_at = datetime.fromisoformat(cached_data['cached_at'])
    ttl = self.cache_strategies[cache_type]['ttl']
    age = (datetime.utcnow() - cached_at).total_seconds()

    # Calculate refresh threshold based on access patterns
    access_frequency = await self.get_access_frequency(key)
    refresh_threshold = ttl * (0.8 - (access_frequency * 0.2)) # 60-80% of TTL

    return age > refresh_threshold

async def warm_popular_cache_entries(self):
    """
    Predictively warm cache for popular queries
    """

    # Get popular query patterns from analytics
    popular_queries = await self.predictor.predict_popular_queries()
```

```
for query_type, queries in popular_queries.items():
    for query in queries:
        cache_key = self.generate_cache_key(query_type, query)

        # Check if already cached
        if not await self.redis.exists(cache_key):
            # Warm cache in background
            asyncio.create_task(self.warm_cache_entry(query_type, query))

async def optimize_cache_size(self):
    """
    Optimize cache size using ML-predicted access patterns
    """

    # Get current cache statistics
    cache_stats = await self.get_cache_statistics()

    # Predict future access patterns
    access_predictions = await self.predictor.predict_access_patterns()

    # Identify candidates for eviction
    eviction_candidates = []
    for key, stats in cache_stats.items():
        if stats['access_frequency'] < access_predictions['threshold']:
            eviction_candidates.append(key)

    # Evict low-value entries
    for key in eviction_candidates[:100]: # Limit batch size
        await self.redis.delete(key)

def generate_cache_key(self, prefix: str, data: Any) -> str:
    """
    Generate consistent cache keys with collision resistance
    """
```

```

    """
    if isinstance(data, dict):
        # Sort keys for consistent hashing
        sorted_data = json.dumps(data, sort_keys=True)
    else:
        sorted_data = str(data)

    data_hash = hashlib.sha256(sorted_data.encode()).hexdigest()[:16]
    return f"{prefix}:{data_hash}"

# Cache warming strategies
CACHE_WARMING_STRATEGIES = {
    'time_based': {
        'schedule': '0 2 * * *', # Daily at 2 AM
        'targets': ['popular_mcc_codes', 'common_risk_factors'],
        'concurrency': 10
    },
    'event_driven': {
        'triggers': ['new_customer_signup', 'bulk_import_completed'],
        'immediate_warm': ['classification_models', 'risk_thresholds']
    },
    'predictive': {
        'model': 'access_pattern_predictor',
        'refresh_interval': '1 hour',
        'confidence_threshold': 0.7
    }
}

```

7.2 Database Performance Optimization

PostgreSQL Performance Tuning

sql

```
-- Database configuration for optimal performance
-- postgresql.conf optimizations

-- Memory settings (for 32GB RAM server)
shared_buffers = '8GB'          -- 25% of RAM
work_mem = '256MB'            -- Per-operation memory
maintenance_work_mem = '2GB'    -- For VACUUM, CREATE INDEX
effective_cache_size = '24GB'   -- 75% of RAM

-- Connection settings
max_connections = 200
superuser_reserved_connections = 3

-- Write-ahead logging
wal_buffers = '64MB'
checkpoint_timeout = '15min'
checkpoint_completion_target = 0.9
wal_keep_size = '2GB'

-- Query planner
random_page_cost = 1.1          -- SSD optimization
effective_io_concurrency = 200    -- SSD concurrent I/O

-- Background writer
bgwriter_delay = '200ms'
bgwriter_lru_maxpages = 100
bgwriter_lru_multiplier = 2.0

-- Auto vacuum settings
autovacuum = on
autovacuum_max_workers = 3
autovacuum_naptime = '20s'
autovacuum_vacuum_threshold = 50
```

```
autovacuum_analyze_threshold = 50

-- Performance monitoring
shared_preload_libraries = 'pg_stat_statements'
track_activity_query_size = 2048
pg_stat_statements.max = 10000
pg_stat_statements.track = all
```

Index Optimization Strategy

```
sql
```

```
-- Core business entity indexes
CREATE INDEX CONCURRENTLY idx_businesses_tenant_lookup
ON businesses(tenant_id, created_at DESC, id)
WHERE deleted_at IS NULL;

CREATE INDEX CONCURRENTLY idx_businesses_search_text
ON businesses USING gin(to_tsvector('english', legal_name || ' ' || COALESCE(dba_name, '')))
WHERE deleted_at IS NULL;

CREATE INDEX CONCURRENTLY idx_businesses_tax_id_hash
ON businesses USING hash(tax_id)
WHERE tax_id IS NOT NULL AND deleted_at IS NULL;

-- Risk assessment performance indexes
CREATE INDEX CONCURRENTLY idx_risk_assessments_recent
ON risk_assessments(business_id, assessed_at DESC, id)
WHERE assessed_at > NOW() - INTERVAL '90 days';

CREATE INDEX CONCURRENTLY idx_risk_assessments_score_range
ON risk_assessments(overall_score, risk_level, assessed_at)
WHERE assessed_at > NOW() - INTERVAL '30 days';

-- Classification lookup indexes
CREATE INDEX CONCURRENTLY idx_classifications_code_lookup
ON business_classifications(classification_type, code, confidence_score DESC)
WHERE is_primary = true;

-- Audit trail partitioned indexes
CREATE INDEX CONCURRENTLY idx_audit_logs_tenant_time
ON audit_logs(tenant_id, created_at DESC)
WHERE created_at > NOW() - INTERVAL '1 year';

-- API usage analytics indexes
```

```
CREATE INDEX CONCURRENTLY idx_api_usage_tenant_endpoint_time
ON api_usage(tenant_id, endpoint, created_at)
WHERE created_at > NOW() - INTERVAL '30 days';

-- Partial indexes for common queries
CREATE INDEX CONCURRENTLY idx_businesses_high_risk
ON businesses(id, legal_name, created_at)
WHERE id IN (
    SELECT business_id
    FROM risk_assessments
    WHERE overall_score > 70
    AND assessed_at > NOW() - INTERVAL '30 days'
);
```

Query Optimization Examples

```
sql
```

```
-- Optimized business search query with full-text search
WITH business_search AS (
    SELECT
        b.id,
        b.legal_name,
        b.dba_name,
        b.website_url,
        b.created_at,
        ts_rank(to_tsvector('english', b.legal_name || ' ' || COALESCE(b.dba_name, '')),
            plainto_tsquery('english', $2)) as relevance_score
    FROM businesses b
    WHERE
        b.tenant_id = $1
        AND b.deleted_at IS NULL
        AND (
            to_tsvector('english', b.legal_name || ' ' || COALESCE(b.dba_name, '')) @@ plainto_tsquery('english', $2)
            OR b.tax_id = $3
        )
    ORDER BY relevance_score DESC, b.created_at DESC
    LIMIT 50
),
risk_data AS (
    SELECT DISTINCT ON (ra.business_id)
        ra.business_id,
        ra.overall_score,
        ra.risk_level,
        ra.assessed_at
    FROM risk_assessments ra
    WHERE ra.business_id IN (SELECT id FROM business_search)
    ORDER BY ra.business_id, ra.assessed_at DESC
)
SELECT
```

```
bs.*,
rd.overall_score,
rd.risk_level,
rd.assessed_at as last_risk_assessment
FROM business_search bs
LEFT JOIN risk_data rd ON bs.id = rd.business_id
ORDER BY bs.relevance_score DESC, bs.created_at DESC;
```

-- Optimized risk trend analysis query

```
WITH monthly_risk_trends AS (
    SELECT
        business_id,
        DATE_TRUNC('month', assessed_at) as month,
        AVG(overall_score) as avg_risk_score,
        COUNT(*) as assessment_count,
        FIRST_VALUE(overall_score) OVER (
            PARTITION BY business_id, DATE_TRUNC('month', assessed_at)
            ORDER BY assessed_at ASC
        ) as month_start_score,
        FIRST_VALUE(overall_score) OVER (
            PARTITION BY business_id, DATE_TRUNC('month', assessed_at)
            ORDER BY assessed_at DESC
        ) as month_end_score
    FROM risk_assessments
    WHERE tenant_id = $1
    AND assessed_at >= NOW() - INTERVAL '12 months'
    GROUP BY business_id, DATE_TRUNC('month', assessed_at), overall_score, assessed_at
)
SELECT
    business_id,
    month,
    avg_risk_score,
    assessment_count,
```

```
month_end_score - month_start_score as risk_change,  
CASE  
    WHEN month_end_score > month_start_score + 5 THEN 'increasing'  
    WHEN month_end_score < month_start_score - 5 THEN 'decreasing'  
    ELSE 'stable'  
END as risk_trend  
FROM monthly_risk_trends  
ORDER BY business_id, month DESC;
```

8. Infrastructure and Deployment Architecture

8.1 Kubernetes Cluster Architecture

Multi-Environment Cluster Design

yaml

```
# Kubernetes cluster configuration
apiVersion: v1
kind: Namespace
metadata:
  name: kyb-production
  labels:
    environment: production
    compliance: "soc2-pci"

---
# Resource quotas for production namespace
apiVersion: v1
kind: ResourceQuota
metadata:
  name: production-quota
  namespace: kyb-production
spec:
  hard:
    requests.cpu: "50"
    requests.memory: "100Gi"
    limits.cpu: "100"
    limits.memory: "200Gi"
    persistentvolumeclaims: "20"
    services: "20"
    secrets: "50"

---
# Network policies for security
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-all
  namespace: kyb-production
```

```
spec:  
  podSelector: {}  
  policyTypes:  
    - Ingress  
    - Egress  
  
---  
# Service mesh configuration (Istio)  
apiVersion: install.istio.io/v1alpha1  
kind: IstioOperator  
metadata:  
  name: kyb-production-istio  
spec:  
  values:  
    global:  
      meshID: kyb-mesh  
      network: kyb-network  
  components:  
    pilot:  
      k8s:  
        resources:  
          requests:  
            cpu: 500m  
            memory: 2048Mi  
      ingressGateways:  
        - name: istio-ingressgateway  
          enabled: true  
          k8s:  
            resources:  
              requests:  
                cpu: 100m  
                memory: 128Mi  
      hpaSpec:
```

```
minReplicas: 2  
maxReplicas: 10
```

Service Deployment Configurations

yaml

```
# API Gateway deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-gateway
  namespace: kyb-production
spec:
  replicas: 3
  selector:
    matchLabels:
      app: api-gateway
  template:
    metadata:
      labels:
        app: api-gateway
        version: v1
    spec:
      containers:
        - name: api-gateway
          image: kyb/api-gateway:1.0.0
          ports:
            - containerPort: 8080
          env:
            - name: REDIS_URL
              valueFrom:
                secretKeyRef:
                  name: redis-credentials
                  key: url
            - name: DB_URL
              valueFrom:
                secretKeyRef:
                  name: postgres-credentials
                  key: url
```

```
resources:
  requests:
    cpu: 200m
    memory: 256Mi
  limits:
    cpu: 500m
    memory: 512Mi
livenessProbe:
  httpGet:
    path: /health
    port: 8080
  initialDelaySeconds: 30
  periodSeconds: 10
readinessProbe:
  httpGet:
    path: /ready
    port: 8080
  initialDelaySeconds: 5
  periodSeconds: 5
```

```
---
# Classification service deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: classification-service
  namespace: kyb-production
spec:
  replicas: 5
  selector:
    matchLabels:
      app: classification-service
  template:
```

```
metadata:  
  labels:  
    app: classification-service  
    version: v1  
  spec:  
    containers:  
      - name: classification-service  
        image: kyb/classification-service:1.0.0  
        ports:  
          - containerPort: 8000  
        env:  
          - name: MODEL_PATH  
            value: "/models"  
          - name: REDIS_URL  
            valueFrom:  
              secretKeyRef:  
                name: redis-credentials  
                key: url  
    resources:  
      requests:  
        cpu: 1000m  
        memory: 2Gi  
      limits:  
        cpu: 2000m  
        memory: 4Gi  
    volumeMounts:  
      - name: model-storage  
        mountPath: /models  
    volumes:  
      - name: model-storage  
        persistentVolumeClaim:  
          claimName: ml-models-pvc
```

```
---
```

```
# Horizontal Pod Autoscaler for classification service
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: classification-service-hpa
  namespace: kyb-production
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: classification-service
  minReplicas: 3
  maxReplicas: 20
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
    - type: Resource
      resource:
        name: memory
        target:
          type: Utilization
          averageUtilization: 80
  behavior:
    scaleDown:
      stabilizationWindowSeconds: 300
    policies:
      - type: Percent
        value: 10
```

```
    periodSeconds: 60
    scaleUp:
      stabilizationWindowSeconds: 60
      policies:
        - type: Percent
          value: 50
        periodSeconds: 60
```

8.2 Multi-Region Deployment Strategy

Global Infrastructure Layout

Global Infrastructure

Primary Region: US-East-1 (Virginia)

- └── Production Kubernetes Cluster (3 AZs)
- └── PostgreSQL Primary Cluster (Multi-AZ)
- └── Redis Cluster (3 nodes across AZs)
- └── Elasticsearch Cluster (3 masters, 6 data nodes)
- └── ML Model Storage (EFS with backup to S3)

Secondary Region: US-West-2 (Oregon)

- └── Disaster Recovery Kubernetes Cluster
- └── PostgreSQL Read Replica
- └── Redis Replica Cluster
- └── Elasticsearch Cross-Cluster Replication
- └── ML Model Sync (S3 Cross-Region Replication)

International Region: EU-West-1 (Ireland)

- └── EU Data Residency Cluster
- └── PostgreSQL EU Cluster (GDPR Compliant)
- └── Redis EU Cluster
- └── Local ML Model Cache
- └── EU-Specific Compliance Services

Edge Locations (Cloudflare CDN)

- └── 200+ Global Edge Locations
- └── Static Asset Caching
- └── API Response Caching (non-sensitive)
- └── DDoS Protection
- └── Bot Mitigation

Cross-Region Data Synchronization

python

```
import asyncio
from dataclasses import dataclass
from typing import List, Dict, Optional
from datetime import datetime

@dataclass
class ReplicationConfig:
    source_region: str
    target_regions: List[str]
    replication_lag_sla: int # seconds
    consistency_level: str # 'eventual' or 'strong'
    encryption_in_transit: bool = True

class MultiRegionDataManager:
    """
    Manages data synchronization across multiple regions
    """

    def __init__(self):
        self.regions = {
            'us-east-1': {'primary': True, 'db_endpoint': 'prod-db-us-east-1'},
            'us-west-2': {'primary': False, 'db_endpoint': 'replica-db-us-west-2'},
            'eu-west-1': {'primary': False, 'db_endpoint': 'eu-db-eu-west-1'}
        }
        self.replication_configs = {
            'businesses': ReplicationConfig(
                source_region='us-east-1',
                target_regions=['us-west-2', 'eu-west-1'],
                replication_lag_sla=30,
                consistency_level='eventual'
            ),
            'risk_assessments': ReplicationConfig(
                source_region='us-east-1',
                target_regions=['us-west-2', 'eu-west-1'],
                replication_lag_sla=60,
                consistency_level='strong'
            )
        }
```

```
        target_regions=['us-west-2'],
        replication_lag_sla=60,
        consistency_level='eventual'
    ),
    'audit_logs': ReplicationConfig(
        source_region='us-east-1',
        target_regions=['us-west-2', 'eu-west-1'],
        replication_lag_sla=300,
        consistency_level='eventual'
    )
}

async def sync_data_to_regions(self, table_name: str, data_changes: List[Dict]) -> Dict:
    """
    Synchronize data changes to target regions
    """

    config = self.replication_configs.get(table_name)
    if not config:
        return {'status': 'error', 'message': f'No replication config for {table_name}'}

    sync_results = {}

    # Parallel sync to all target regions
    sync_tasks = []
    for region in config.target_regions:
        task = self.sync_to_region(region, table_name, data_changes)
        sync_tasks.append(task)

    # Wait for all syncs to complete
    results = await asyncio.gather(*sync_tasks, return_exceptions=True)

    for i, result in enumerate(results):
        region = config.target_regions[i]
```

```
if isinstance(result, Exception):
    sync_results[region] = {'status': 'error', 'error': str(result)}
else:
    sync_results[region] = result

return sync_results

async def handle_region_failover(self, failed_region: str) -> Dict:
    """
    Handle failover when a region becomes unavailable
    """

    failover_plan = {
        'us-east-1': {
            'primary_failover': 'us-west-2',
            'traffic_routing': 'dns_failover',
            'data_consistency_check': True
        },
        'us-west-2': {
            'primary_failover': 'us-east-1',
            'traffic_routing': 'load_balancer',
            'data_consistency_check': False
        }
    }

    plan = failover_plan.get(failed_region)
    if not plan:
        return {'status': 'error', 'message': f'No failover plan for {failed_region}'}

    # Execute failover procedures
    failover_steps = [
        self.update_dns_routing(failed_region, plan['primary_failover']),
        self.promote_read_replica(plan['primary_failover']),
        self.update_application_config(plan['primary_failover']),
    ]
```

```
        self.verify_service_health(plan['primary_failover'])
    ]

    if plan['data_consistency_check']:
        failover_steps.append(self.verify_data_consistency())

    results = await asyncio.gather(*failover_steps)

    return {
        'status': 'completed',
        'failed_region': failed_region,
        'new_primary': plan['primary_failover'],
        'failover_time': datetime.utcnow(),
        'steps_completed': len([r for r in results if r.get('status') == 'success'])
    }
```

8.3 CI/CD Pipeline Architecture

GitOps Deployment Pipeline

yaml

```
# GitHub Actions CI/CD Pipeline
name: KYB Platform CI/CD

on:
  push:
    branches: [main, develop, 'feature/*']
  pull_request:
    branches: [main, develop]

env:
  REGISTRY: ghcr.io
  IMAGE_NAME: kyb-platform

jobs:
  # Security and code quality
  security-scan:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3

      - name: Run Trivy vulnerability scanner
        uses: aquasecurity/trivy-action@master
        with:
          scan-type: 'fs'
          scan-ref: '!'
          format: 'sarif'
          output: 'trivy-results.sarif'

      - name: Upload Trivy scan results
        uses: github/codeql-action/upload-sarif@v2
        with:
          sarif_file: 'trivy-results.sarif'
```

```
- name: Run Semgrep security scan
  uses: returntocorp/semgrep-action@v1
  with:
    config: auto
    publishToken: ${{ secrets.SEMGREP_APP_TOKEN }}

# Build and test services
build-and-test:
  runs-on: ubuntu-latest
  needs: security-scan
  strategy:
    matrix:
      service: [api-gateway, classification-service, risk-service, data-ingestion]

  steps:
    - uses: actions/checkout@v3

    - name: Set up Go
      if: matrix.service == 'api-gateway' || matrix.service == 'data-ingestion'
      uses: actions/setup-go@v3
      with:
        go-version: 1.21

    - name: Set up Python
      if: matrix.service == 'classification-service' || matrix.service == 'risk-service'
      uses: actions/setup-python@v4
      with:
        python-version: '3.11'

    - name: Cache dependencies
      uses: actions/cache@v3
      with:
        path: |
```

```
~/.cache/go-build
~/go/pkg/mod
~/.cache/pip
key: ${{ runner.os }}-${{ matrix.service }}-${{ hashFiles('**/go.mod', '**/requirements.txt') }}

- name: Install dependencies and run tests
  run: |
    cd services/${{ matrix.service }}
    if [ -f "go.mod" ]; then
      go mod download
      go test -v -race -coverprofile=coverage.out ./...
      go tool cover -html=coverage.out -o coverage.html
    elif [ -f "requirements.txt" ]; then
      pip install -r requirements.txt
      python -m pytest --cov=. --cov-report=html --cov-report=xml
    fi

- name: Upload coverage reports
  uses: codecov/codecov-action@v3
  with:
    file: ./coverage.xml
    flags: ${{ matrix.service }}

# Build container images
build-images:
  runs-on: ubuntu-latest
  needs: build-and-test
  if: github.ref == 'refs/heads/main' || github.ref == 'refs/heads/develop'

  strategy:
    matrix:
      service: [api-gateway, classification-service, risk-service, data-ingestion]
```

```
steps:
- uses: actions/checkout@v3

- name: Set up Docker Buildx
  uses: docker/setup-buildx-action@v2

- name: Log in to Container Registry
  uses: docker/login-action@v2
  with:
    registry: ${{ env.REGISTRY }}
    username: ${{ github.actor }}
    password: ${{ secrets.GITHUB_TOKEN }}

- name: Extract metadata
  id: meta
  uses: docker/metadata-action@v4
  with:
    images: ${{ env.REGISTRY }}/{{ env.IMAGE_NAME }}-${{ matrix.service }}
    tags: |
      type=ref,event=branch
      type=ref,event=pr
      type=sha,prefix={{branch}}-

- name: Build and push Docker image
  uses: docker/build-push-action@v4
  with:
    context: ./services/${{ matrix.service }}
    push: true
    tags: ${{ steps.meta.outputs.tags }}
    labels: ${{ steps.meta.outputs.labels }}
    cache-from: type=gha
    cache-to: type=gha,mode=max
```

```
# Deploy to staging

deploy-staging:
  runs-on: ubuntu-latest
  needs: build-images
  if: github.ref == 'refs/heads/develop'
  environment: staging

  steps:
    - uses: actions/checkout@v3

    - name: Setup Kubectl
      uses: azure/setup-kubectl@v3
      with:
        version: 'v1.28.0'

    - name: Setup Helm
      uses: azure/setup-helm@v3
      with:
        version: 'v3.12.0'

    - name: Configure kubectl
      run: |
        echo "${{ secrets.KUBE_CONFIG_STAGING }}" | base64 -d > ~/.kube/config
        kubectl config use-context staging

    - name: Deploy with Helm
      run: |
        helm upgrade --install kyb-platform ./deploy/helm/kyb-platform \
          --namespace kyb-staging \
          --create-namespace \
          --values ./deploy/helm/values-staging.yaml \
          --set image.tag=${{ github.sha }} \
          --wait --timeout=600s
```

```
- name: Run smoke tests
  run: |
    kubectl wait --for=condition=ready pod -l app=api-gateway -n kyb-staging --timeout=300s
    # Run basic API health checks
    kubectl exec -n kyb-staging deployment/api-gateway -- curl -f http://localhost:8080/health

# Deploy to production
deploy-production:
  runs-on: ubuntu-latest
  needs: build-images
  if: github.ref == 'refs/heads/main'
  environment: production

steps:
- uses: actions/checkout@v3

- name: Deploy to production
  run: |
    # GitOps - commit to argocd repository
    git clone https://${{ secrets.GITOPS_TOKEN }}@github.com/org/kyb-gitops.git
    cd kyb-gitops

# Update image tags in production manifests
yq eval '.spec.template.spec.containers[0].image = "${{ env.REGISTRY }}/{{ env.IMAGE_NAME }}-api-' yq eval '.spec.template.spec.containers[0].image = "${{ env.REGISTRY }}/{{ env.IMAGE_NAME }}-clas'

# Commit and push changes
git config user.name "GitHub Actions"
git config user.email "actions@github.com"
git add .
git commit -m "Deploy KYB Platform ${{ github.sha }} to production"
git push origin main
```

9. Monitoring and Observability

9.1 Comprehensive Monitoring Stack

Three Pillars of Observability

```
yaml
```

```
# Prometheus configuration for metrics collection

global:
  scrape_interval: 15s
  evaluation_interval: 15s
  external_labels:
    cluster: 'kyb-production'
    replica: '1'

rule_files:
  - "kyb-alerts.yml"
  - "sla-alerts.yml"
  - "business-metrics.yml"

scrape_configs:
  # Kubernetes components
  - job_name: 'kubernetes-nodes'
    kubernetes_sd_configs:
      - role: node
    relabel_configs:
      - source_labels: [__address__]
        regex: '(:\d+):10250'
        target_label: __address__
        replacement: '${1}:9100'
      - action: labelmap
        regex: __meta_kubernetes_node_label_(.+)

  # Application services
  - job_name: 'kyb-services'
    kubernetes_sd_configs:
      - role: endpoints
    namespaces:
      names: ['kyb-production']
    relabel_configs:
```

```
- source_labels: [__meta_kubernetes_service_annotation_prometheus_io_scrape]
  action: keep
  regex: true
- source_labels: [__meta_kubernetes_service_annotation_prometheus_io_path]
  action: replace
  target_label: __metrics_path__
  regex: (.+)

# Business metrics (custom application metrics)
- job_name: 'business-metrics'
  static_configs:
    - targets: ['api-gateway:8080', 'classification-service:8000']
  metrics_path: '/metrics/business'
  scrape_interval: 30s

# Database metrics
- job_name: 'postgres-exporter'
  static_configs:
    - targets: ['postgres-exporter:9187']
  scrape_interval: 30s

# Redis metrics
- job_name: 'redis-exporter'
  static_configs:
    - targets: ['redis-exporter:9121']
  scrape_interval: 15s

alerting:
  alertmanagers:
    - static_configs:
        - targets: ['alertmanager:9093']
```

Custom Business Metrics

python

```
from prometheus_client import Counter, Histogram, Gauge, CollectorRegistry
import time

# Business-specific metrics
BUSINESS_METRICS_REGISTRY = CollectorRegistry()

# API usage metrics
api_requests_total = Counter(
    'kyb_api_requests_total',
    'Total API requests by endpoint and tenant',
    ['endpoint', 'method', 'status_code', 'tenant_id'],
    registry=BUSINESS_METRICS_REGISTRY
)

api_request_duration = Histogram(
    'kyb_api_request_duration_seconds',
    'API request duration in seconds',
    ['endpoint', 'method'],
    registry=BUSINESS_METRICS_REGISTRY,
    buckets=[0.1, 0.25, 0.5, 1.0, 2.0, 5.0, 10.0]
)

# Business operation metrics
business_verifications_total = Counter(
    'kyb_business_verifications_total',
    'Total business verifications processed',
    ['result', 'tenant_id', 'verification_type'],
    registry=BUSINESS_METRICS_REGISTRY
)

risk_assessments_total = Counter(
    'kyb_risk_assessments_total',
    'Total risk assessments performed',
    ['tenant_id', 'risk_type']
)
```

```
[risk_level', 'tenant_id', 'assessment_type'],
registry=BUSINESS_METRICS_REGISTRY
)

classification_accuracy = Gauge(
    'kyb_classification_accuracy_ratio',
    'Current classification model accuracy',
    ['model_version', 'classification_type'],
    registry=BUSINESS_METRICS_REGISTRY
)

# System health metrics

active_tenants = Gauge(
    'kyb_active_tenants_total',
    'Number of active tenants',
    registry=BUSINESS_METRICS_REGISTRY
)

cache_hit_rate = Gauge(
    'kyb_cache_hit_rate_ratio',
    'Cache hit rate by cache type',
    ['cache_type'],
    registry=BUSINESS_METRICS_REGISTRY
)

ml_model_inference_time = Histogram(
    'kyb_ml_model_inference_duration_seconds',
    'ML model inference time',
    ['model_name', 'model_version'],
    registry=BUSINESS_METRICS_REGISTRY,
    buckets=[0.01, 0.05, 0.1, 0.25, 0.5, 1.0, 2.0]
)
```

```
# Revenue and business metrics
subscription_revenue_gauge = Gauge(
    'kyb_monthly_recurring_revenue_dollars',
    'Monthly recurring revenue in USD',
    ['plan_type'],
    registry=BUSINESS_METRICS_REGISTRY
)

customer_churn_rate = Gauge(
    'kyb_customer_churn_rate_ratio',
    'Monthly customer churn rate',
    registry=BUSINESS_METRICS_REGISTRY
)

class MetricsCollector:
    """Collect and expose custom business metrics"""

    def __init__(self):
        self.registry = BUSINESS_METRICS_REGISTRY

    def record_api_request(self, endpoint: str, method: str, status_code: int,
                          duration: float, tenant_id: str):
        """Record API request metrics"""
        api_requests_total.labels(
            endpoint=endpoint,
            method=method,
            status_code=status_code,
            tenant_id=tenant_id
        ).inc()

        api_request_duration.labels(
            endpoint=endpoint,
            method=method
```

```
    ).observe(duration)

    def record_business_verification(self, result: str, tenant_id: str,
                                      verification_type: str):
        """Record business verification completion"""
        business_verifications_total.labels(
            result=result,
            tenant_id=tenant_id,
            verification_type=verification_type
        ).inc()

    def update_model_accuracy(self, model_version: str, classification_type: str,
                             accuracy: float):
        """Update ML model accuracy metrics"""
        classification_accuracy.labels(
            model_version=model_version,
            classification_type=classification_type
        ).set(accuracy)

    def update_business_metrics(self):
        """Update business KPI metrics (called periodically)"""
        # This would typically fetch from database/analytics system
        pass
```

SLA and Alerting Rules

yaml

```
# kyb-alerts.yml - Prometheus alerting rules

groups:
- name: kyb-sla-alerts

rules:

# API availability SLA (99.99%)
- alert: APIAvailabilityBelowSLA
  expr: |
    (
      rate(kyb_api_requests_total{status_code!="5.."}[5m]) /
      rate(kyb_api_requests_total[5m])
    ) < 0.9999
  for: 1m
  labels:
    severity: critical
    sla: availability
  annotations:
    summary: "API availability below 99.99% SLA"
    description: "API availability is {{ $value | humanizePercentage }} over the last 5 minutes"

# API latency SLA (95th percentile < 2 seconds)
- alert: APILatencyAboveSLA
  expr: |
    histogram_quantile(0.95,
      rate(kyb_api_request_duration_seconds_bucket[5m])
    ) > 2.0
  for: 2m
  labels:
    severity: warning
    sla: latency
  annotations:
    summary: "API latency above 2 second SLA"
    description: "95th percentile latency is {{ $value }}s over the last 5 minutes"
```

```
# Classification accuracy below threshold
- alert: ClassificationAccuracyLow
  expr: kyb_classification_accuracy_ratio < 0.95
  for: 5m
  labels:
    severity: warning
    component: ml-model
  annotations:
    summary: "Classification accuracy below 95%"
    description: "Model {{ $labels.model_version }} accuracy is {{ $value | humanizePercentage }}"

# High error rate
- alert: HighErrorRate
  expr: |
    rate(kyb_api_requests_total{status_code=~"5.."}[5m]) /
    rate(kyb_api_requests_total[5m]) > 0.01
  for: 1m
  labels:
    severity: critical
    component: api
  annotations:
    summary: "High error rate detected"
    description: "Error rate is {{ $value | humanizePercentage }} over the last 5 minutes"

- name: kyb-business-alerts
  rules:

# Revenue impact alerts
- alert: SignificantRevenueDroppage
  expr: |
    (
      kyb_monthly_recurring_revenue_dollars -
```

```

kyb_monthly_recurring_revenue_dollars offset 24h
) / kyb_monthly_recurring_revenue_dollars offset 24h < -0.1
for: 1h
labels:
severity: critical
impact: revenue
annotations:
summary: "MRR dropped by more than 10%"
description: "Monthly recurring revenue decreased by {{ $value | humanizePercentage }}"

# Tenant health monitoring
- alert: TenantInactive
expr: |
  time() - max by (tenant_id) (kyb_api_requests_total) > 86400
labels:
severity: warning
component: customer-success
annotations:
summary: "Tenant {{ $labels.tenant_id }} inactive for 24+ hours"
description: "No API activity detected for tenant {{ $labels.tenant_id }}"

```

This completes the comprehensive Technical Architecture Document covering all aspects of system design, performance optimization, infrastructure, and operations. The document provides detailed implementation guidance for:

- 1. Complete data architecture** with PostgreSQL, Redis, and Elasticsearch
- 2. AI/ML model architectures** using in-house/open-source models
- 3. External API integrations** prioritizing free/open-source options
- 4. Security and compliance** frameworks for SOC 2, PCI DSS, and GDPR
- 5. Performance optimization** strategies and caching
- 6. Multi-region infrastructure** and deployment architecture

7. Comprehensive monitoring and observability stack

The architecture is designed to support the complete KYB platform from MVP through enterprise scale, with specific focus on cost optimization, security, and performance requirements outlined in the executive overview.

Would you like me to proceed with creating the **Feature Specifications Document** next?